

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/337549406>

# Data Scraping, Ingestion, and Modeling: Bringing Data from cars.com into the Intro Stats Class

Article in CHANCE · October 2019

DOI: 10.1080/09332480.2019.1695443

CITATIONS

0

READS

7

2 authors, including:



Nicholas Jon Horton

Amherst College

229 PUBLICATIONS 9,906 CITATIONS

SEE PROFILE

Some of the authors of this publication are also working on these related projects:



Eating Disorders [View project](#)



Helping students to "think with data" [View project](#)

# Data scraping, ingestion, and modeling: bringing data from cars.com into the intro stats class

*Sarah McDonald and Nicholas J. Horton*

*August 28, 2018*

## Introduction

New tools have made it much easier for students to develop skills to work with interesting data sets as they begin to extract meaning from data. To fully appreciate the statistical analysis cycle, students benefit from repeated experiences collecting, ingesting, wrangling, analyzing data and communicating results. How can we bring such opportunities into the classroom? We describe a classroom activity, originally developed by Danny Kaplan (Macalester College), in which students can expand upon statistical problem solving by hand-scraping data from cars.com, ingesting these data into R, then carrying out analyses of the relationships between price, mileage, and model year for a selected type of car.

Most students might be interested in car prices since many will be purchasing a car at some point in the near future. This activity can help them develop better understanding of factors associated with car prices.

The revised GAISE (Guidelines for Assessment and Instruction in Statistics Education) College report (2016) notes the importance of multivariate thinking and the use of technology. Car prices, model year, and mileage are all factors to consider when purchasing or selling a car. Introductory statistics courses need to move beyond only addressing bivariate questions to be able to explore multivariate relationships.

In an increasingly data-rich society, plenty of information is available to prospective car purchasers. Consumers

can analyze and compare multiple cars to try to get the best deal. By gathering data by hand from cars.com then using this information to generate multivariable visualizations and model prices, students gain experience (1) working in groups, (2) practicing undertaking reproducible analyses, and (3) exploring a multivariate dataset. These key ideas of data generation, data ingestion, data visualization for multivariate analyses, and data modeling are reinforced throughout the activity.

We begin by describing the activity, sharing examples of data, visualizations, and models, then suggesting possible extensions and providing concluding thoughts. Instructor materials and datasets associated with this activity can be found at <https://github.com/Amherst-Statistics/Cars-Scraping-Webinar>.

## **Activity: Class One**

Students work in pairs of two and use two computers to gather and hand-enter data concerning the cost of a specific model of a car, then analyze the variations in pricing, price associations with mileage and age, the rate at which cars depreciate, and the cost of driving one mile. One student reads off data from cars.com and the other enters the data into a spreadsheet. Each pair is assigned a different city.

The first step of the activity involves gathering data from *cars.com*. Using the ‘advanced filter’ option, the model and make of the car are specified, along with the assigned location and restriction to recent years. Various components in price determination include the model, year, mileage, and location.

As an example, Figure 1 features a 2015 Toyota Prius from the Dallas area, priced at \$17,998 with 15,866 miles whereas the 2014 Toyota Prius is priced lower at \$10,995 but with a higher mileage of 81,076.

30 miles ▾

75219

New/Used

☐ All
☐ New
☒ Used

☐ Certified Pre-Owned

Year

2015 ▾

2011 ▾


Make

☒ Toyota
☐ Acura
☐ BMW
☐ Chevrolet
☐ Ford
☐ Honda

[Show all Makes in A-Z](#)

Model

☒ Prius



USED

\$17,998

15,866 mi.

2015 Toyota Prius Three

GOOD DEAL

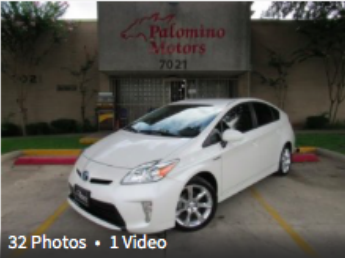
HOT CAR

Exterior Color: Barcelona Red Metallic • Interior Color: Dark Gray • Transmission: 1-Speed CVT w/OD • Drivetrain: FWD

24 Photos • 1 Video

C

Free CARFAX 1-Owner Report



USED

\$10,995

81,076 mi.

2014 Toyota Prius Two

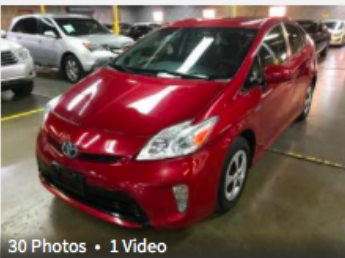
GOOD DEAL

Exterior Color: Blizzard Pearl • Interior Color: Dark Gray • Transmission: Automatic • Drivetrain: Front Wheel Drive

32 Photos • 1 Video

C

Free CARFAX 1-Owner Report



USED

\$7,999

137,331 mi.

2013 Toyota Prius 5dr HB One

GREAT DEAL

Exterior Color: Red • Interior Color: Gray • Transmission: A • Drivetrain: Front Wheel Drive

30 Photos • 1 Video

C

Free CARFAX 1-Owner Report

Figure 1: Cars.com example snapshot of used Toyota Prius models from Dallas

|     | A      | B     | C | D     | E    | F     | G      |
|-----|--------|-------|---|-------|------|-------|--------|
| 280 | Toyota | Prius | 2 | 16711 | 2013 | 28394 | Dallas |
| 281 | Toyota | Prius | 5 | 16552 | 2013 | 38390 | Dallas |
| 282 | Toyota | Prius | 2 | 16480 | 2013 | 32528 | Dallas |
| 283 | Toyota | Prius | 3 | 15280 | 2012 | 39585 | Dallas |
| 284 | Toyota | Prius | 2 | 14991 | 2014 | 50290 | Dallas |
| 285 | Toyota | Prius | 2 | 14988 | 2014 | 28462 | Dallas |
| 286 | Toyota | Prius | 2 | 14982 | 2013 | 22836 | Dallas |
| 287 | Toyota | Prius | 3 | 14980 | 2013 | 46670 | Dallas |
| 288 | Toyota | Prius | 4 | 14854 | 2013 | 44106 | Dallas |
| 289 | Toyota | Prius | 2 | 14781 | 2015 | 43521 | Dallas |
| 290 | Toyota | Prius | 3 | 14599 | 2012 | 45487 | Dallas |

Figure 2: Student hand-scraped data for Dallas entered into an Excel spreadsheet

Figure 2 illustrates data gathered and entered into an Excel sheet for a group assigned to find car prices in Dallas.

The data are entered into a spreadsheet (e.g., Excel, Open Office, or Google Spreadsheet) using a template `cars.csv` to ensure that the variable names are consistent between groups. Once the group has completed the hand-scraping of 30 or 35 cars, they will upload this spreadsheet into RStudio and run an instructor provided RMarkdown file (`cars.Rmd`). The RMarkdown file reads the data that they have uploaded to generate descriptive statistics, creates multivariate displays, and fits a multiple regression model. The students need to interpret the results and add their descriptions into the file.

The scatterplot produced in Figure 3 uses student-gathered data for Toyota Prius to display the relationship between prices and mileage for Dallas cars. The scatterplot reflects how car prices depreciate as a function of mileage and model year. After the car's first year, the discrepancy in price based on mileage by year tends to diminish.

The plot below displays a linear regression model for Prius prices in Dallas.

Here the `ggformula` interface to the `ggplot2` graphics system is used because it provides a general modeling syntax similar to the `lm()` function in R.




Figure 3: Toyota Prius prices in Dallas based on mileage

```
library(ggformula)

gf_point(price ~ mileage, color = ~ year, data = Dallas) %>%

  gf_lm()
```

|             | Estimate   | Std. Error | t value | Pr(> t ) |
|-------------|------------|------------|---------|----------|
| (Intercept) | 19721.0125 | 706.8204   | 27.90   | 0.0000   |
| mileage     | -0.1075    | 0.0135     | -7.96   | 0.0000   |

The students then edit the RMarkdown file to interpret their results based on the model and the graphical displays. For the Dallas group, the summary output of the model in the table suggests that for every mile driven, the car's predicted value (determined by price) will decrease on average by about eleven cents.

Common errors that students experience include issues with formatting (e.g., if they included dollar signs in the column for price) or problems where they used different variable names than specified in the assignment.

To obtain credit for the first part of the assignment, students must:

- 1) post the formatted file to RPubS (to allow a brief discussion of student findings and interpretations)
- 2) email the csv file to the instructor

## Activity: Class Two

Prior to the next class period, the instructor collates the data from each group (in csv files) to create graphical displays, multiple regression models, and interpretations from the data from all of the cities. These results can be referenced as part of a future class discussion. The collation process will identify issues (e.g., inconsistent formatting or variable naming) in the individual datasets, which also provide an opportunity for discussion.

Figure 4 displays the scatterplot visualizing the relationship between the price and mileage, where an interaction is included between the mileage and (categorical) model year, using data scraped from all of the cities ( $n = 830$ ).

```
library(mosaic)

tally(~ location, data = ds)
```

```
## location

##      40202      Atlanta  Bangor, ME  Baton Rouge      Boston
##      40      40      40      40      40
##      Buffalo      Chicago      Cleveland      Dallas      Los Angeles
##      33      41      26      41      40
##      Minneapolis      New Orleans      NYC      Phoenix      Portland
##      59      33      40      39      40
##      Richmond Salt Lake City      San Diego      San Francisco      Seattle
##      40      33      39      39      39
##      Tampa
##      40
```

We note that one group has included the zip code (needed to specify location in cars.com) instead of the city name. Also note that some groups only scraped 33 or 39 cars (to keep the class together on day one data scraping was cut off after a certain amount of time).

```
gf_point(price ~ mileage, color = ~ year, data = ds) %>%
  gf_lm() %>%
  gf_labs(y = "price (US $)")
```

The multiple regression output describes the relationship between the price based on location, mileage, year, and the interaction between mileage and year. This is a relatively sophisticated model, with 32 predictors. Example interpretations of this model are included below:

LOCATION: After controlling for mileage and year, prices for a Toyota Prius in Boston are predicted to be \$564 less than in Louisville, Kentucky (the reference group). (Note the reference group is the first group in






Figure 4: Price versus mileage stratified by year ( $n = 830$ )

|                        | Estimate   | Std. Error | t value | Pr(> t ) |
|------------------------|------------|------------|---------|----------|
| (Intercept)            | 17061.0694 | 868.8562   | 19.64   | 0.0000   |
| locationAtlanta        | -1638.4149 | 462.5576   | -3.54   | 0.0004   |
| locationBangor, ME     | -1689.6974 | 463.9047   | -3.64   | 0.0003   |
| locationBaton Rouge    | -745.2125  | 474.3208   | -1.57   | 0.1166   |
| locationBoston         | -563.6481  | 460.0693   | -1.23   | 0.2209   |
| locationBuffalo        | -581.6074  | 484.2352   | -1.20   | 0.2301   |
| locationChicago        | -2237.4990 | 456.4975   | -4.90   | 0.0000   |
| locationCleveland      | -1491.5866 | 520.8768   | -2.86   | 0.0043   |
| locationDallas         | -1078.1113 | 462.0475   | -2.33   | 0.0199   |
| locationLos Angeles    | 2319.6793  | 460.0475   | 5.04    | 0.0000   |
| locationMinneapolis    | -622.8958  | 423.7223   | -1.47   | 0.1419   |
| locationNew Orleans    | -573.2974  | 498.8439   | -1.15   | 0.2508   |
| locationNYC            | -594.5619  | 458.8934   | -1.30   | 0.1955   |
| locationPhoenix        | -325.9632  | 463.8124   | -0.70   | 0.4824   |
| locationPortland       | 65.2454    | 461.6683   | 0.14    | 0.8876   |
| locationRichmond       | -744.3217  | 461.1860   | -1.61   | 0.1069   |
| locationSalt Lake City | -1954.0469 | 494.6800   | -3.95   | 0.0001   |
| locationSan Diego      | 257.6979   | 461.9773   | 0.56    | 0.5771   |
| locationSan Francisco  | 1578.2819  | 461.3929   | 3.42    | 0.0007   |
| locationSeattle        | 2136.5419  | 463.0608   | 4.61    | 0.0000   |
| locationTampa          | -2152.2974 | 462.1671   | -4.66   | 0.0000   |
| mileage                | -0.0606    | 0.0095     | -6.38   | 0.0000   |
| year2012               | -251.3108  | 1135.1085  | -0.22   | 0.8248   |
| year2013               | 3237.2317  | 894.6854   | 3.62    | 0.0003   |
| year2014               | 3140.1907  | 888.3434   | 3.53    | 0.0004   |
| year2015               | 3252.5139  | 885.3063   | 3.67    | 0.0003   |
| year2016               | 8208.6105  | 874.4768   | 9.39    | 0.0000   |
| mileage:year2012       | 0.0171     | 0.0144     | 1.19    | 0.2345   |
| mileage:year2013       | -0.0180    | 0.0121     | -1.48   | 0.1394   |
| mileage:year2014       | -0.0034    | 0.0140     | -0.25   | 0.8060   |
| mileage:year2015       | -0.0099    | 0.0139     | -0.71   | 0.4778   |
| mileage:year2016       | -0.1819    | 0.0275     | -6.60   | 0.0000   |

the data set, which by R's default is alphabetically. Here, it is Louisville, Kentucky as one group entered location as a zip code, 40202, rather than by name.)

MILEAGE: Holding location constant, the predicted price of a Prius decreases on average by about six cents for an additional mile for Priuses of the model.

INTERACTION: The interaction of mileage and year is more complicated to interpret, since it includes five regression coefficients. We would predict an additional average decrease in value of about eighteen cents per mile driven for 2016 models compared with 2011 models, after accounting for location. This is a great example of the *new car effect*: there is a much higher rate of depreciation in value of newer cars in comparison to older models.

Other aspects of the model lend themselves to discussion. There are two outliers (both from the same group) with very low prices. These are likely prices that were entered incorrectly. In addition, the functional form of the relationship between price and mileage (conditional on year) is not very linear (though the regression model is assuming linear relationships). We consider these as part of possible extensions of the activity.

## Extensions

In terms of introductory statistics, this activity works to develop students ability to undertake the entire data analysis cycle. They collect data by scraping information (by hand) from a website, then loading this into RStudio.

With the data set, students can practice interpreting interaction terms in the model. This practice will prove beneficial to students as data sets (and models) become increasingly complex in future statistics courses.

In the model produced in Figure 4, two outliers are observed. The two points can be found in the data set by searching for Toyota Priuses priced well below the average. Both data points indicate a pricing at \$2,500 from Chicago, with one 2014 model and one 2015 model, and both of the same model type (four). The 2014 model has a mileage of 17,152 wherein the average price for a used car of similar mileage in Chicago is around \$15,550 and the 2015 model (with current mileage of 21,027) would be priced around \$16,000, according to the model. It appears that the large discrepancy between the price and mileage (well under the average predicted price by \$13,000) could be due to input error, such as a missing zero at the end of the value. Students should note these outliers and decide from inference whether or not to include them in the final model.

We have introduced this activity early in the course so have not focused much on the functional form of the relationship between price and mileage (beyond noting that the relationship is not very linear, see Figure 5). Consideration of more flexible regression models could be undertaken to better reflect the underlying relationships.

While students included additional information in their spreadsheets regarding trim models or add-on packages




Figure 5: Price versus mileage for 2013 vehicles (with superimposed smoother)

for the cars, this was not incorporated into the modeling. Additional data wrangling would be needed to bring this into the model as an additional predictor given the inconsistent and idiosyncratic ways that such information is made available by sellers in *cars.com*.

Potential pitfalls include that the predictions made from the linear models reflect only the cars in the data set and are not completely representative of all car prices and locations. The models produced also do not reflect consumer habits in its entirety as the data gathered only demonstrates cars that are for sale and not necessarily sales price: negotiation is important in determining sales price! Aspects of these biases and data limitations could form the basis of a discussion of design.

## Conclusions

This activity is intended to reinforce critical aspects outlined by the GAISE report, including teamwork, problem solving, and the use of data to make decisions. This activity encourages multivariate thinking through application facilitated by technology. The discovery of the *new car effect* is not obvious in a bivariate analysis.

Additional concepts such as data ingestion, regression modeling, and graphical visualizations are among the other key learning outcomes.

Students are given the opportunity to gather data by hand and build models to extract meaningful inferences. The learning objectives of the cars.com activity permeate through other spheres of consumer habits and students gain independence in their ability to make the best consumer decisions. Financial literacy is an important capacity for students to develop. This activity may help prepare students to make better decisions when buying a car.

A focus on conceptual understanding, integration of real data with a context and purpose, and a fostering of active learning are also critical to students' comprehension. The usage of technology to explore concepts and and analyze data, and assessments to improve and evaluate student learning are additional goals of this

activity.

## Further Reading

GAISE College Report ASA Revision Committee, *Guidelines for Assessment and Instruction in Statistics Education College Report 2016*, <http://www.amstat.org/education/gaise>.

National Academies of Sciences, Engineering, and Medicine. 2018. *Data Science for Undergraduates: Opportunities and Options*. Washington, DC: The National Academies Press. <https://doi.org/10.17226/25104>, <https://nas.edu/envisioningds>

Ben Baumer, Mine Cetinkaya-Rundel, Andrew Bray, Linda Loi, & Nicholas J. Horton (2014). RMarkdown: Integrating A Reproducible Analysis Tool into Introductory Statistics. *Technology Innovations in Statistics Education*, 8(1). Retrieved from <https://escholarship.org/uc/item/90b2f5xh>

Randall Pruim, Daniel T. Kaplan, and Nicholas J. Horton. “The Mosaic Package: Helping Students to Think with Data Using R.” *R*, June 2017, [journal.r-project.org/archive/2017/RJ-2017-024/](http://journal.r-project.org/archive/2017/RJ-2017-024/).

Nicholas J. Horton, Benjamin S. Baumer, & Hadley Wickham (2015) Taking a Chance in the Classroom: Setting the Stage for Data Science: Integration of Data Management Skills in Introductory and Second Courses in Statistics, *CHANCE*, 28:2, 40-50, DOI: 10.1080/09332480.2015.1042739

## Biographies

Sarah McDonald is a student at Amherst College, majoring in Statistics. Her areas of interest include applications of statistical analysis in consumer purchasing and behavioral habits. Her undergraduate research involves studying effective ways to integrate and facilitate computation in introductory statistics courses.

Nicholas J. Horton is Beitzel Professor of Technology and Society and Professor of Statistics and Data Science at Amherst College, with interests in longitudinal regression, missing data methods, statistical computing,

and statistical education. He received his doctorate in biostatistics from the Harvard School of Public Health in 1999, and has co-authored a series of books on statistical computing and data science.